

EVERYWHERE.

Redis Data Structures for Non-Redis people

Alvin Richards
Chief Education Officer
Redis Labs



#Redis #RedisConf

presented by:
redislabs
home of redis

Today's ramble

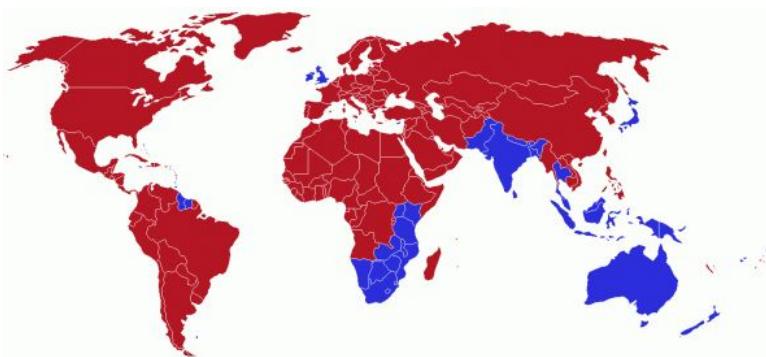
- Data Structures in Redis
- Transactions in Redis
- How to think & approach Redis
- Conclusions

Everything is Different... in Britain

I really ate these things



I drove on this side of the road



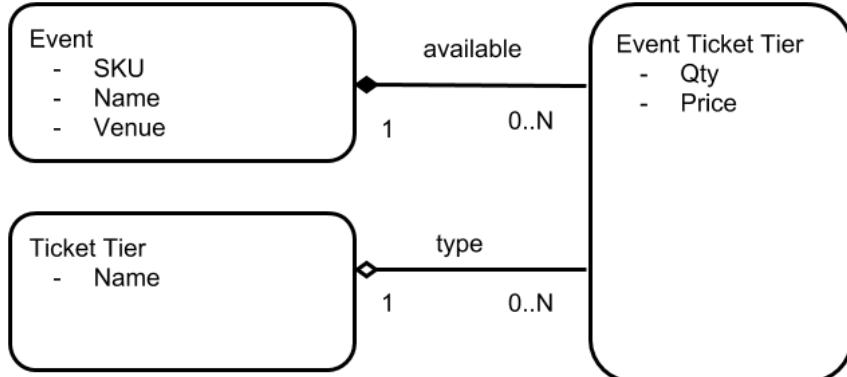
I learnt to "translate"



And you probably grew up with these...

Entity-Relation Diagrams

Tables/Foreign keys



XML

```
<?xml version="1.0"?>
<event_catalog>
  <event>
    <sku>123-ABC-723</sku>
    <name>Men's 100m Final</name>
    <disabled_access>true</disabled_access>
    <medal_event>true</medal_event>
    <venue>Olympic Stadium</venue>
    <category>Track & Field</category>
    <capacity>60102</capacity>
    <available>
      <general>
        <qty>20000</qty>
        <price>25.00</price>
      </general>
    </available>
  </event>
</event_catalog>
```

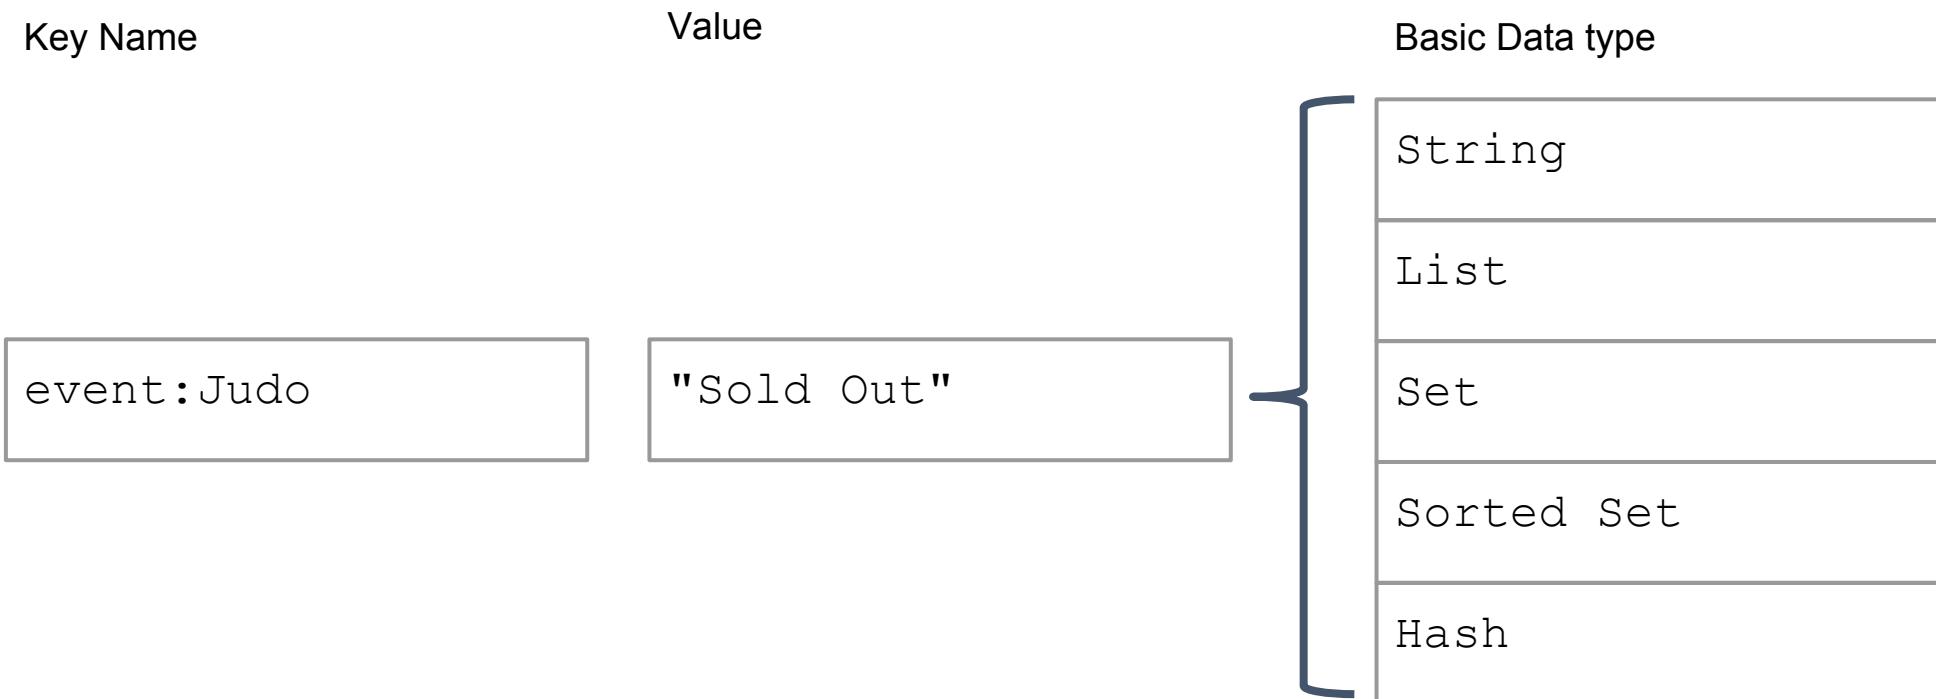
JSON

```
{  
    'sku': "123-ABC-723",  
    'name': "Men's 100m Final",  
    'disabled_access': True,  
    'medal_event': True,  
    'venue': "Olympic Stadium",  
    'category': "Track & Field",  
    'capacity': 60102,  
    'available':  
        { 'general': { 'qty': 20000,  
                      'price': 25.00  
                    }  
    }  
}
```

Redis is different... and here's how

- Redis is a Multi Model Key store
- Commands operate on Keys
 - Commands work on specific data types
 - Commands work on specific encoding of values
- Data types of Keys can change overtime

Keys, Values & Data types



Keys

Key Name

```
event:Judo
```

Keys

- Unique
- Binary Safe
- Up to 512MB in size

Key Naming

Key Name

```
event:Judo
```

```
event:venues
```

```
event:Fencing
```

```
event:fencing
```

```
Event:fencing
```

Naming

- Up to you
- Be consistent
- Colon is a typical community convention
 - Domain Object Name
 - Colon
 - Unique ID
- Plurals are nice for Lists, Sets
- Case-sensitive

Key Expiration

Key Name

Expiration

```
> set event:Fencing ex 50
```

event:Fencing

50

event:Judo

```
> pexpire event:Judo 2
```

event:Judo

2

Expiration or TTL

- Default - keys are retained
- Specified in
 - Seconds
 - Milliseconds
 - Unix Epoch
- Added to Key
- Removed from Key

String Datatype & Encoding

Key Name	Value	Encoding
> set event:Judo "Sold Out"	event:Judo "Sold Out"	embstr
> set event:Judo 42	event:Judo "42"	int
> set event:Judo 3.14159	event:Judo "3.14159"	embstr

Command Execution, Data types and Encoding

Key Name	Value	Data type	Encoding	
> incr event:Fencing	event:Fencing 42	String	int	✓
> incr event:Judo	event:Judo "Sold Out"	String	embstr	✗
> incr events	events Fencing, Judo	List		✗

Lists

Key Name

```
events
```

Lists

- Ordered collection of Strings
- Duplicates allowed
- Elements added
 - Left, Right, by position

Left



```
> lpop events  
"Judo"
```

```
> rpop events  
"Taekwondo"
```

Why are Redis Lists cool?

Useful to create

- Queues, Stacks
- Capped Lists

Typical use cases

- Interprocess communications

Many operations are $O(1)$

- Same time complexity regardless of List size
- LPOP, LPUSH, LPUSHX
- RPOP, RPUSH, RUSHPX
- BLPOP, BRPOP
- RPOPLPUSH, BRPOPLPUSH
- LLEN

Sets

Key Name

events

Member/Element

"Judo"

"Taekwondo"

"Fencing"

```
> smembers events
1) "Judo"
2) "Taekwondo"
3) "Fencing"
```

Sets

- Unordered collection of unique Strings
- Set operations
 - Difference
 - Intersect
 - Union

Why are Redis Sets cool?

Set Operations

- Difference, Intersect, Union

Typical use cases

- Tag Cloud
- Unique Visitors

Many operations are O(1)

- Same time complexity regardless of Set size
- SADD
- SCARD
- SPOP, SRANDMEMBER

Sorted Sets

Key Name

events

Score Value

2	"Judo"
3	"Taekwondo"
1	"Fencing"

```
> zrange events 0 1
1) "Fencing"
2) "Judo"
```

Sets

- Ordered collection of unique Strings
- Manipulation by value, score, position or lexicography
- Set operations
 - Intersect
 - Union

Why are Redis Sorted Sets cool?

Set Operations

- Intersect & Union

Typical use cases

- Leaderboard
- Priority Queue

Many operations are $O(1)$ or $O(\log N)$

- ZADD,
- ZCOUNT, ZINCRBY
- ZSCORE, ZRANK

Hashes

Key Name	Score	Value
event:Judo	venue	Super Dome
	capacity	32000
	subway	True

```
> hget event:Judo capacity  
"32000"
```

Hashes

- Field & Value pairs
- Single level
- Dynamically add remove fields

Why are Redis Hashes cool?

Set Operations

- Intersect & Union

Typical use cases

- Rate Limiting
- Session Cache

Many operations are O(1)

- HGET, HSET, HSETNX
- HINCRBY, HINCRBYFLOAT
- HLEN, HSTRLEN
- HEXISTS

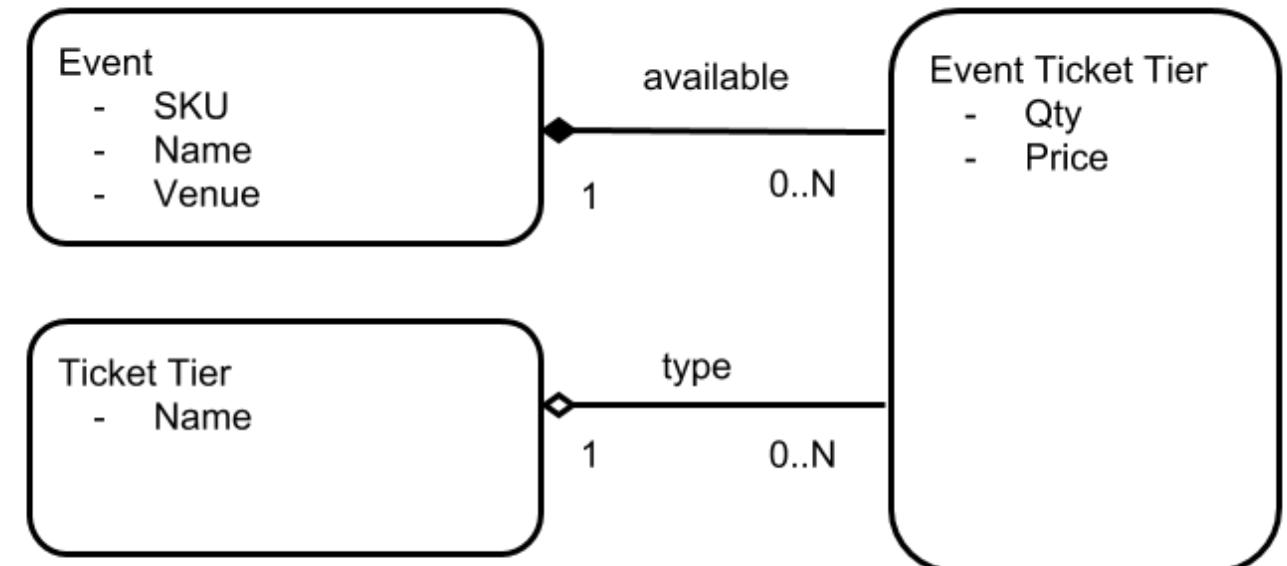
Like... Wow... dude



Putting that into practice

Modeling an Event & Ticket Tiers

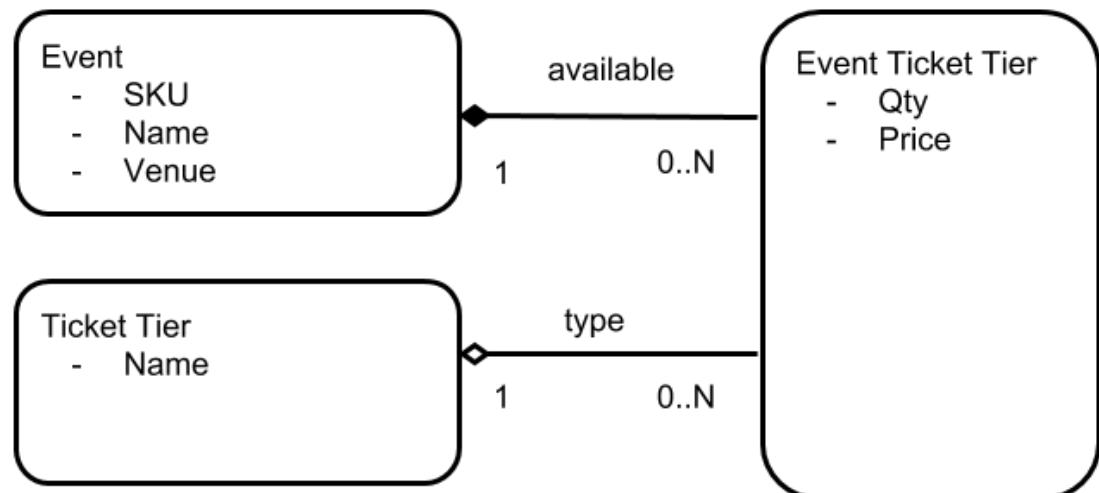
```
{  
  'sku': "123-ABC-723",  
  'name': "Men's 100m Final",  
  'disabled_access': True,  
  'medal_event': True,  
  'venue': "Olympic Stadium",  
  'category': "Track & Field",  
  'capacity': 60102,  
  'available':  
    { 'general': { 'qty': 20000,  
                  'price': 25.00  
                }  
  }  
}
```



Embedding

Relationships

Associations, Cardinality & Containment



- **Association & Cardinality**
 - An Event Ticket Tier is associated with One and Only One Event
 - An Event may have zero or more Event Ticket Tiers
 - An Event Ticket Tier is of One and Only One Ticket Tier
 - A Ticket Tier may have zero or more Event Ticket Tiers
- **Containment**
 - When the Event is removed, Event Ticket Tiers are also removed

Solution #1: Strings

The most common (and worst)

```
> set event:event:123-ABC-723
"{"sku": "123-ABC-723",
"category": "Track & Field",
"available": {"general": {
"price": 25.0, "qty": 20000}},
"capacity": 60102, "name": "Men's 100m Final",
"disabled_access": true,
"medal_event": true, "venue": "Olympic Stadium"}"
```

- Good
 - Simple to get and set the whole object
 - Many frameworks do this by default
- Bad
 - Many frameworks do this by default
 - Resource usage
 - CPU, Memory, Network
 - Serialization / Deserialization
 - Whole object must be Get/Set
 - e.g. Cannot just increment a single field
 - Painful to deal with

Solution #2: Hash

Typical recommendation

```
> hmset event:123-ABC-723
sku "123-ABC-723"
name "Men's 100m Final"
disabled_access True
medal_event True
venue "Olympic Stadium"
category "Track & Field"
capacity 60102
available:general:qty 20000
available:general:price 22.50
```

- Hashes are arbitrarily sized
- Field names up to 512MB
 - But really, should be short!
- Flatten hierarchy / Associations
 - Large number of price/qty pairs can be created
 - Ticket Tiers can be removed by deleting pairs of fields
- Containment
 - When the Key is removed, all fields are also removed

Solution #2: Hash

```
> hmset event:123-ABC-723
sku "123-ABC-723"
name "Men's 100m Final"
disabled_access True
medal_event True
venue "Olympic Stadium"
category "Track & Field"
capacity 60102
available:general:qty 20000
available:general:price 22.50
```

- Good
 - Atomic updates
 - Atomic deletes
- Bad
 - Large objects

Solution #3: Multiple Hashes plus Set Uncommon

```
> hmset event:123-ABC-723
```

```
sku "123-ABC-723"  
name "Men's 100m Final"  
disabled_access True  
medal_event True  
venue "Olympic Stadium"  
category "Track & Field"  
capacity 60102
```

```
> hset event:123-ABC-723:general
```

```
qty 20000  
price 22.50
```

```
> sadd event:123-ABC-723:available
```

```
event:123-ABC-723:general
```

- Hash for each Object
 - Event
 - Tier
- (Optional) Set to hold Relationship
 - Redundant
 - Useful if you need to only detect presence of Relationship end

Solution #3: Multiple Hashes plus Set

```
> hmset event:123-ABC-723
sku "123-ABC-723"
name "Men's 100m Final"
disabled_access True
medal_event True
venue "Olympic Stadium"
category "Track & Field"
capacity 60102
> hset event:123-ABC-723:general
qty 20000
price 22.50
> sadd event:123-ABC-723:available
event:123-ABC-723:general
```

- Good
 - Extensible structures
 - Independently stored
- Bad
 - Relationship has to be manually maintained
 - All Keys need to be on same Shard

We're cool, right?



Transactions

- Redis has Transactions!
 - Atomic
 - Isolated
- Commands queued
- Queued commands executed sequentially as a single unit

```
> multi  
> set event:Judo 100  
> incr event:Judo  
> exec  
  
> multi  
> set event:Judo "Sold Out"  
> discard
```

Transactions in Action #1

Client 1



Client 2

get event:Judo

get event:Judo

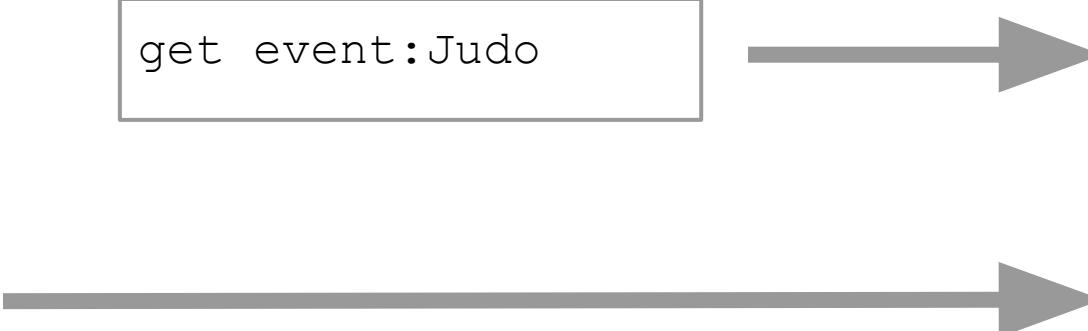
set event:Judo 100

100

100

101

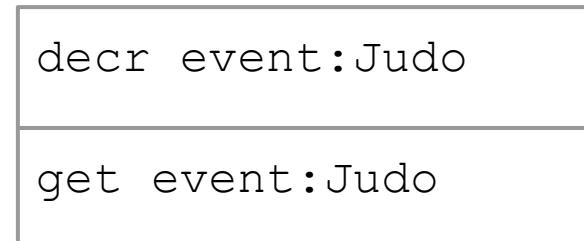
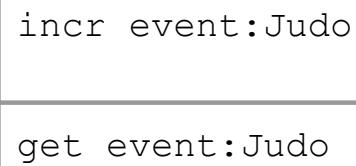
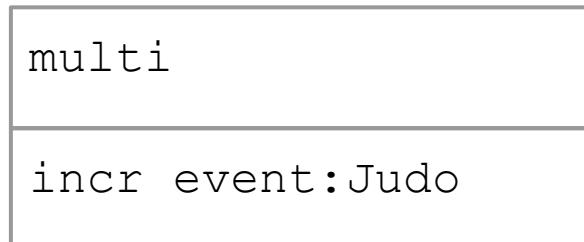
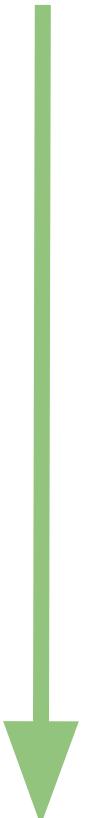
101



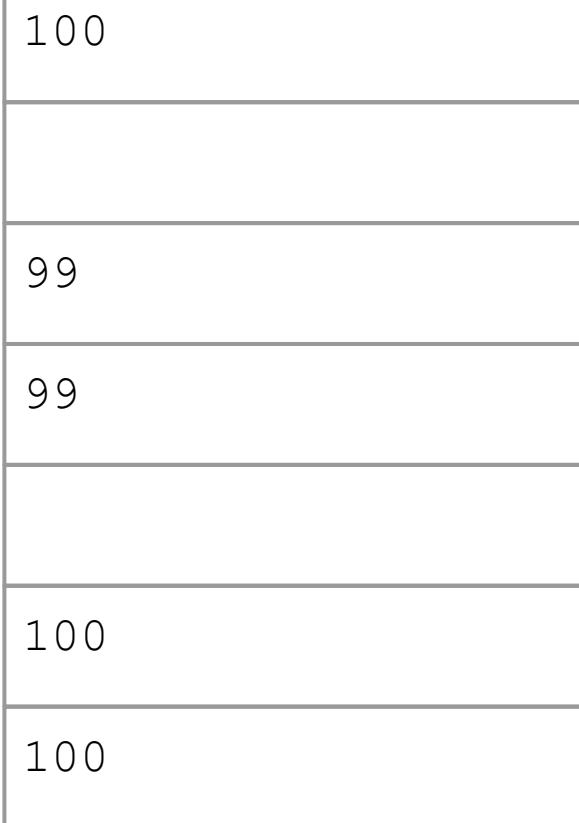
Time

Transactions in Action #2

Client 1



```
set event:Judo 100
```

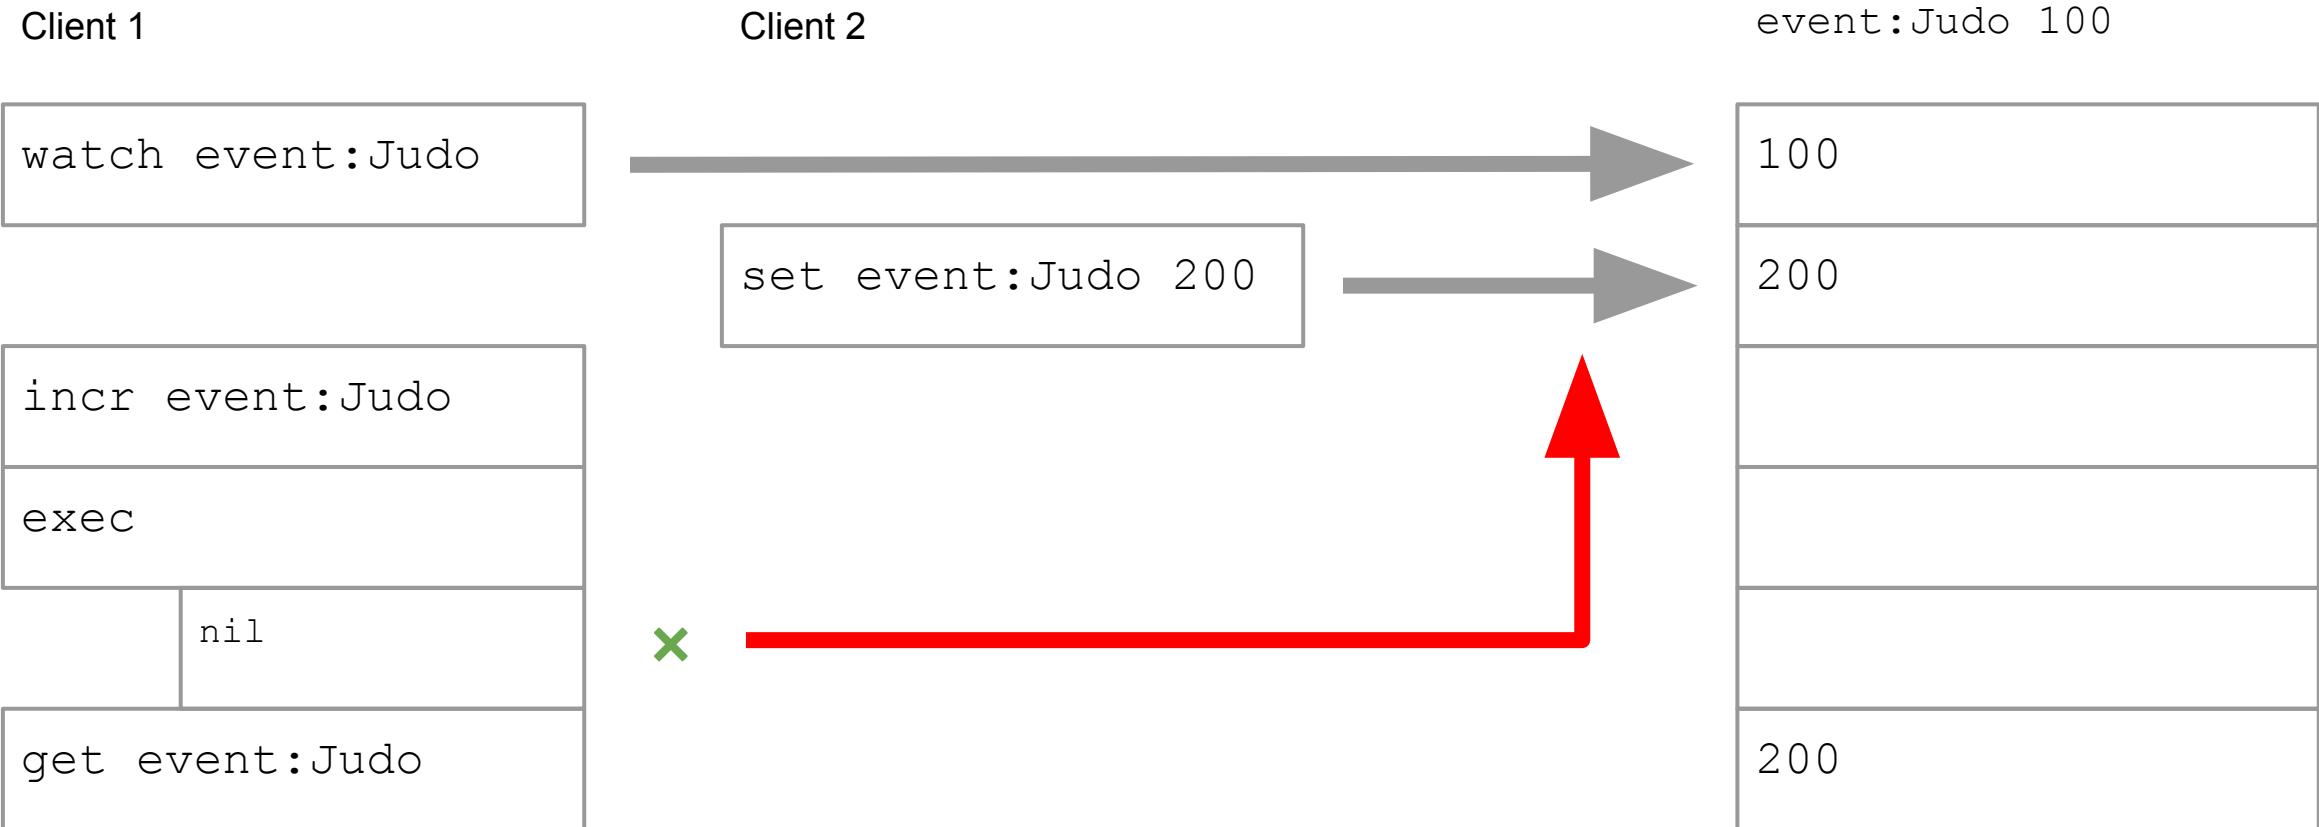


Time

Optimistic Concurrency

- WATCH
- UNWATCH

Optimistic Concurrency in Action #1



Why are Transaction important?

- Allows data to be safely modeled in many Keys
- Watches allow for Optimistic Concurrency Control

Really, are cool now?



Geospatial Add point

```
> geoadd venues 139.75 35.693333  
"Nippon Budokan"
```

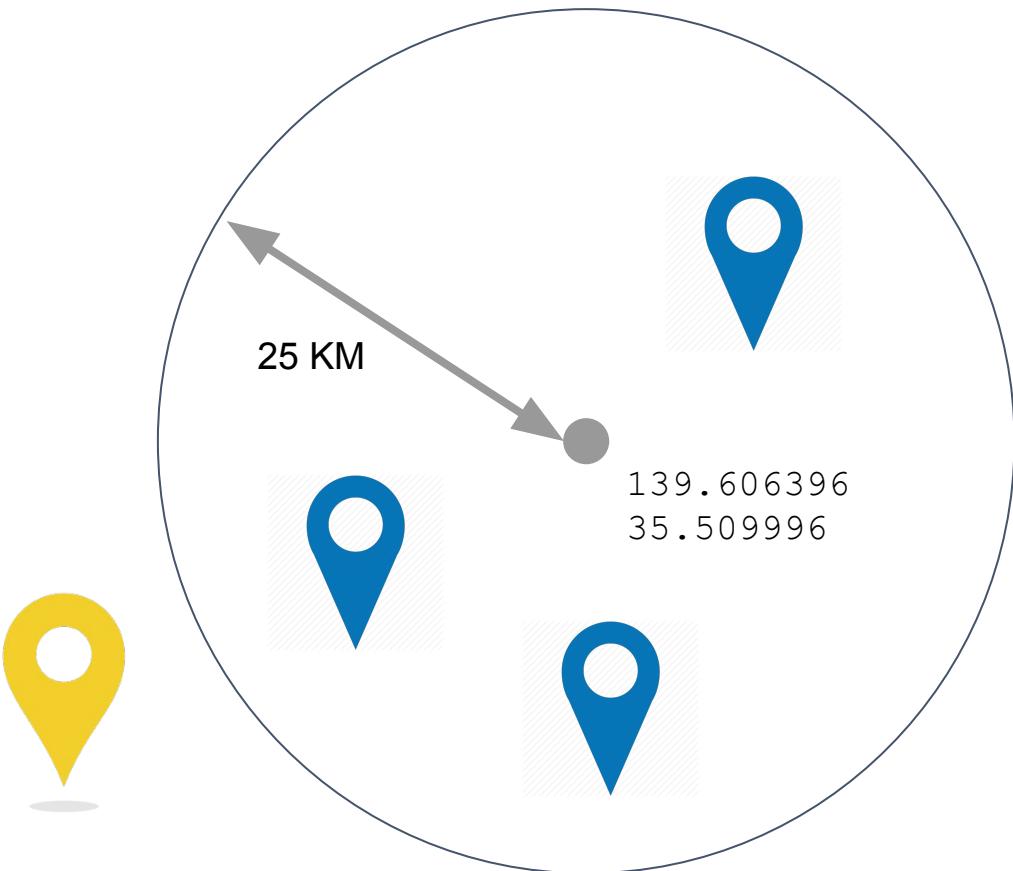
Geospatial

- Named point
- Position
 - Long/Lat coordinates
 - 52 bit Geo hash computed
- Stored in Sorted Set

Geospatial

Within radius of point

```
> georadius venues 139.606396 35.509996 25 km
```



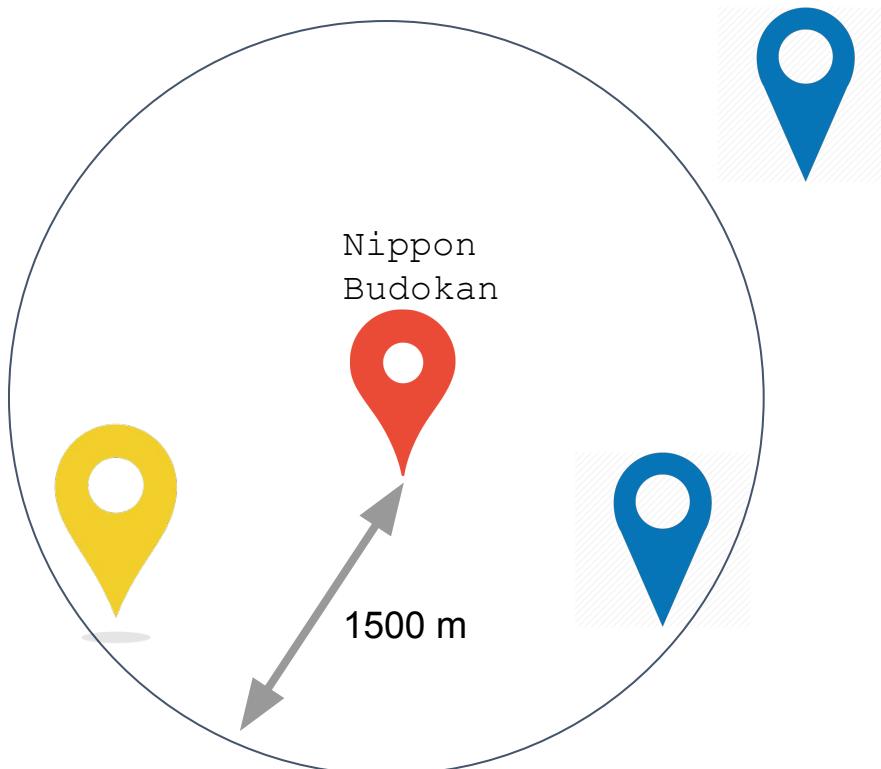
Geo Radius

- Distance from Long/Lat
 - Units Meters, Km, Feet, Miles
- Return
 - Distance from point
 - Coordinates of point
 - Geo Hash

Geospatial

Within radius of Named point

```
> georadiusbymember venues "Nippon Budokan" 1500 m
```

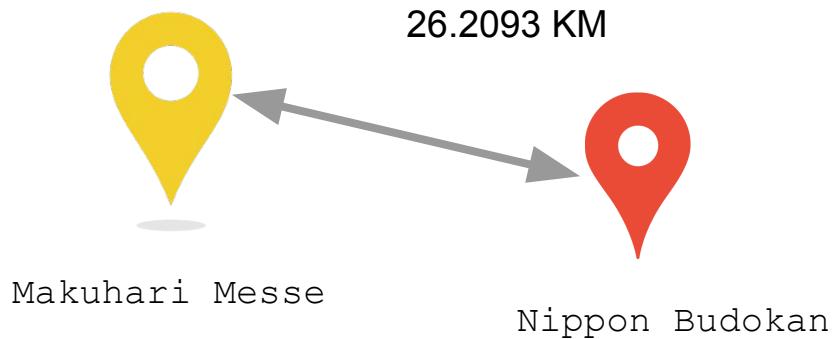


Geo Radius

- Distance from Named point
 - Units Meters, Km, Feet, Miles
- Return
 - Distance from point
 - Coordinates of point
 - Geo hash

Geospatial Distance between named points

```
> geodist venues "Makuhari Messe" "Nippon Budokan" km
```



Geo Distance

- Distance between two points
 - Units Meters, Km, Feet, Miles
- Return
 - Distance from point
 - Coordinates of point

Why Geospatial is cool

Use case

- AdTech
 - Find audience to broadcast offer
 - Find offers near me
- Social
 - Which friends are close
 - Where's a good place to eat near Market & Church?
- Logistics
 - Drop off / Pick ups, routing

Publish / Subscribe

```
> publish news:Judo "Event starts in 1  
hour"
```

```
> subscribe news:Judo
```

```
> psubscribe news:*
```

```
> unsubscribe news:Judo
```

Publish / Subscribe

- Publish to a single channel
- Subscribe
 - One or more channels
 - Glob-style wildcards

Publish / Subscribe

Client 1

```
publish news:Judo  
"Medal event about  
to start"
```

news:Judo

"Medal event about
to start"

```
publish news:Judo  
"Event starting"
```

"Event Starting"

```
publish news:Judo  
"Event finished"
```

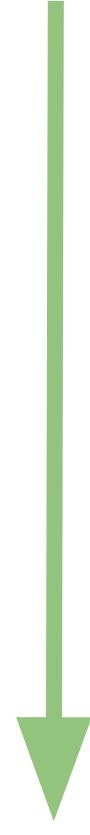
"Event Finished"

Client 2

```
subscribe news:Judo
```

"Event Finished"

Time



Why Publish / Subscribe is cool

Use case

- Interprocess communications

Conclusions

Data is modeled different in Redis

- Understand
 - Your Use Case & operations required
 - Data cardinality & Distribution
 - Select Data Structure
- Think about
 - Flat name space - what is your naming convention?
 - Transactions - atomic, enable safe multi key writes
 - Relationships - flatten into Hashes, or combine with Sets

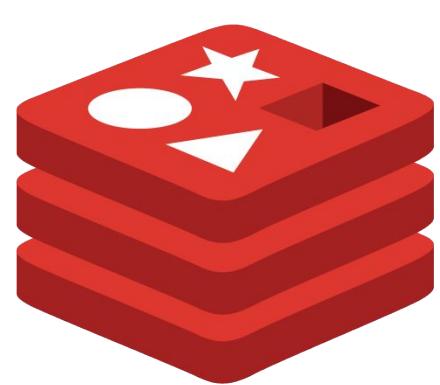
And things we did not talk about

HyperLogLog - probabilistic counting structure

Bitfields - compact bit structures

Modules

- ReJSON - full JSON support
- Redisearch - text search



redis UNIVERSITY

And finally...

- RU101 - An Introduction to Redis Data Structures
- Registration starts 25th April
- Course starts early June
 - Online, classroom style
 - 6 weeks duration
 - 2-3 hours commitment per week
 - Virtual Labs for every student
 - Certificate of Completion



EVERYWHERE.

Thank You



#Redis #RedisConf

presented by:

redislabs
home of redis

Key Namespaces

Key Name

```
event:Judo
```

Key namespace

- Logical Database
 - Default is db[0]
- Flat namespace
- Typically used for separation within an application / domain